

Assignment 2

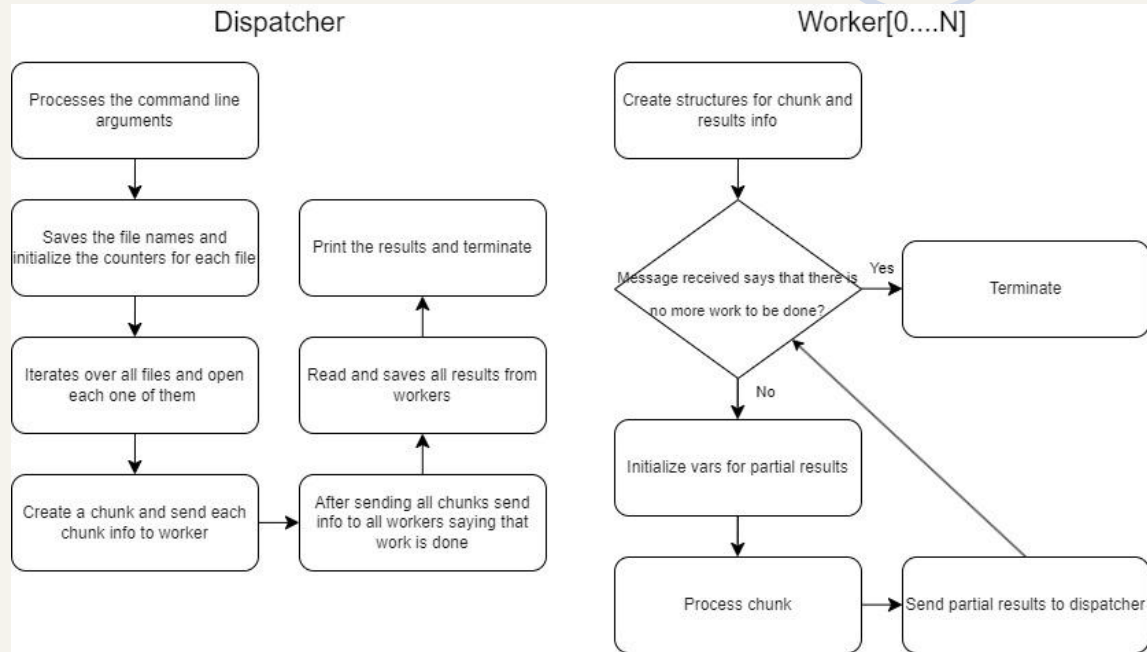
Multiprocessing

CLE - Computação em Larga Escala
Maio 2023

Eduardo Fernandes - 98512
Alexandre Pinto - 98401

Text Processing - Multiprocessing Implementation

Objective: Given a text file count the total number of words, and also the number of words containing an a, e, i, o, u, y.
The multiprocessing implementation splits the work among the worker processes.



Text Processing - Results

Machine used in testing: Laptop LeNovo Legion, Intel Core i7-10750H Hexa Core 12 Threads, Nvidia GeForce RTX 2060, RAM 16GB

The results below were obtained with a chunk size of 4096 and processing all text files.

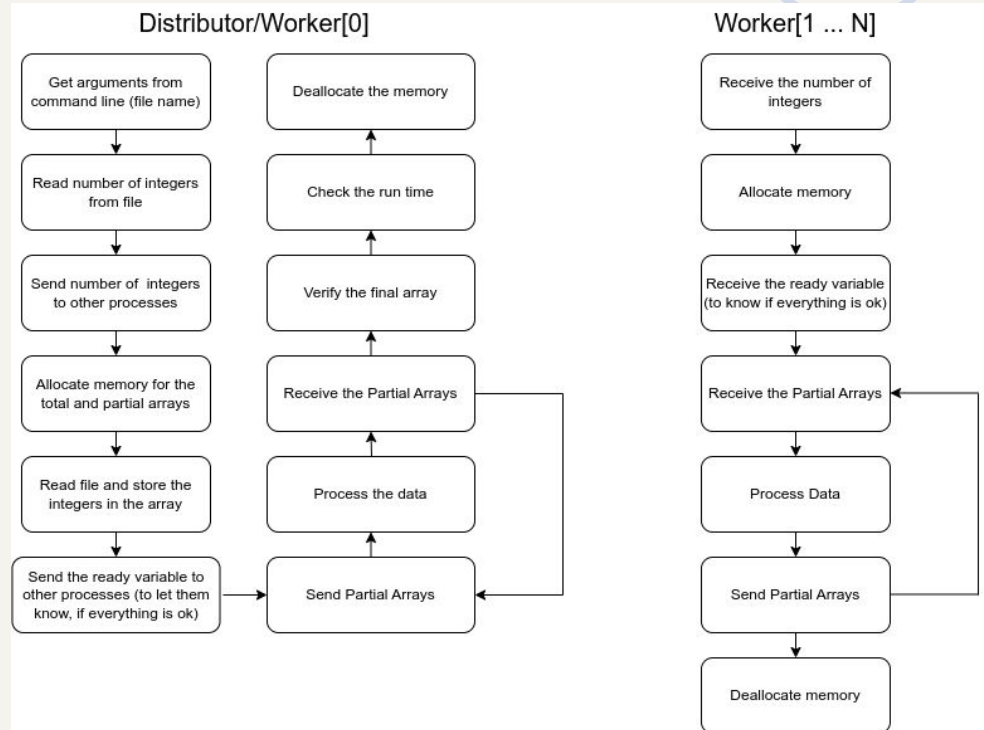
- **1 Worker** (one dispatcher process and one worker process) : Elapsed Time = $1.52 \cdot 10^{-3}$ s;
- **2 Worker** (one dispatcher process and two worker processes) : Elapsed Time = $8.26 \cdot 10^{-4}$ s;
- **4 Worker** (one dispatcher process and four worker processes) : Elapsed Time = $7.52 \cdot 10^{-4}$ s;
- **8 Worker** (one dispatcher process and eight worker processes) : Elapsed Time = $7.20 \cdot 10^{-4}$ s;

Command:

```
mpicc -Wall -O3 -o textProcessing textProcessing.c textProcessingFunctions.c  
mpiexec -n [number_processes] ./textProcessing -f [files to be processed]
```

Integers Sorting - Multiprocessing Implementation

Objective: Given a binary file sort the integers there presented. The multiprocessing implementation uses a distributor process to split the work among the worker processes. However, the distributor is also responsible for some work, it is the worker[0].



Integers Sorting - Results

Machine used in testing: Laptop LeNovo Legion, Intel Core i7-10750H Hexa Core 12 Threads, Nvidia GeForce RTX 2060, RAM 16GB

Objective: Given a binary file sort the integers there presented. The multiprocessing implementation uses a distributor process to split the work among the worker processes. However, the distributor is also responsible for some work, it is the worker[0].

Number Processes \ File Name	1	2	4	8	T I M E (s)
dataSeq32.bin	$5.23 \cdot 10^{-5}$	$8.58 \cdot 10^{-5}$	$9.19 \cdot 10^{-5}$	$1.64 \cdot 10^{-4}$	
dataSeq256K.bin	$8.07 \cdot 10^{-2}$	$6.70 \cdot 10^{-2}$	$6.56 \cdot 10^{-2}$	$8.86 \cdot 10^{-2}$	
dataSeq1M.bin	$3.20 \cdot 10^{-1}$	$2.85 \cdot 10^{-1}$	$2.44 \cdot 10^{-1}$	$3.22 \cdot 10^{-1}$	
dataSeq16M.bin	5.90	5.96	4.59	5.19	

Commands:

```
mpicc -Wall -O3 sorting.c -o sorting -lm  
mpirun -np [number_processes] ./sorting [file_name]
```

Conclusion

After looking at our results, we came to the following conclusions:

- In the program 1, the increase on the number of workers will improve the run time because there are many processes doing the work.
- In the program 2, the run time tends to become faster when using more processes to sort the integers. However, in small files, the communications between processes can use more time than the sorting itself. So, it is faster to use less processes in these cases.
- Finally, the increase on the number of processes should be better when having big files to process.